Renyi Tang

Email: tangrenyi99@gmail.com Mobile: 872-216-7537
LinkedIn Website: renyi.live Address: Pittsburgh, PA

EDUCATION

- **Carnegie Mellon University** — Pittsburgh, PA
  M.S. in Machine Learning. **GPA**: 4.0/4.0 — Aug 2022 – Dec 2023 (expected)

- **University of Chicago** — Chicago, IL
  M.S. in Computer Science; B.S. in Mathematics. **GPA**: 3.82/4.0 — Sep 2017 – Dec 2021
  - Selected courses: DL Systems, Networks and Distributed Systems, Database Systems, Parallel Computing; Advanced ML, Advanced DL, Probabilistic Graphical Models, ML and Text Mining, On-Device ML, Computer Vision

WORK EXPERIENCE

- **LinkedIn** — Sunnyvale, CA
  Software Engineering Intern — May 2023 - now
  - Restructured and optimized Spark Scala dataframe schemas and transformations, saving HDFS storage by 20%.
  - Developed foundational Python code on negative sampling, supporting all online metric learning tasks at LinkedIn.
  - Enhanced CTR prediction metrics at LinkedIn Feed by 1% through sequence modeling on user interaction data.
  - Adopted agile software development principles by collaborating in a team setting and using online tools such as GitHub, JIRA and Confluence to organize workflow.

- **Huawei** — Shenzhen, China
  Software Engineering Intern — Jan 2022 - July 2022
  - Built scalable visualization tools for large-scale point cloud data in Python and improved render speed with C++ and CUDA by 75%. Enabled features not available in the industry's latest open-source packages.
  - Designed and deployed a 3D video data collection system in Python and C++. Enabled synchronization between different data sources through parallel programming to increase throughput and frame match accuracy by 21%.
  - Enhanced deep learning models based on PointNet(++) and PCPNet to improve multi-view stereo quality by 4%.
  - Created data preprocessing pipeline with PCL, Open3D to reduce the memory footprint of training models by 50%.

- **Squirrel AI** — Pittsburgh, PA
  Software Engineering Intern — Jun 2019 - Aug 2019
  - Implemented pipelines to feed and store one million daily pen stroke data with HDFS in Java.
  - Built real-time communication between front-end and back-end through the WebSocket network protocol, enabling live recognition of handwriting. Designed a smooth and aesthetic interface that raised user satisfaction.
  - Integrated high-performance C++ code into Java backend with JNI, saving 30+ days to rebuild the wheel.
  - Researched and packaged an algorithm that improved the accuracy of recognizing simple math equations to 95%.

PROJECTS

- **Needle, a mini PyTorch (Python, C++, CUDA)**
  - Built an auto-differentiation framework with NumPy, CPU, CUDA backends.
  - Supported high-performance matrix computation through tiling and shared memory, reducing training time by 40%.
  - Added advanced features powering the latest research such as subgraph backpropagation and implicit layers.

- **TwitterOps, a lambda architecture big data project (Java, Scala)**
  - Organized 100GB Twitter dataset into tables and analyzed account behavior. Deployed with AWS EC2 and S3.
  - Built batch layer by saving data to HDFS with Hadoop serialization and Hive, serving layer by extracting information with HBase, Spark, Spark SQL, and speed layer with Kafka to send data and maintain queues.

- **SimpleDB, a prototype database (Java)**
  - Implemented key components of a database from scratch, including physical and logical data management, operators, query optimization and transaction management.
  - Wrote comprehensive unit tests to benchmark performance over a variety of contexts.

SKILLS

- Python (advanced), Java, C/C++, CUDA, JavaScript, HTML/CSS, Node.js; PyTorch (advanced), TensorFlow, Spark, Hadoop, Hive, Kafka, Docker, Git, SQL